

Boeing Data Science Challenge

Hazen Eckert

December 7, 2023

I used Python 3 in a Jupyter notebook to do all my work. I chose these based on familiarity and ease of use for this small dataset. For data cleaning and exploration I utilized Pandas. Natural Language Toolkit was used to generate features from strings. Sci-Kit Learn was used for the machine learning models and model validation.

I first cleaned the data. I imputed the values for features that were null for very few entries and added a label for null entries for the other features. A few entries had unique labels for a specific trim. These were relabeled using the most common label for that trim package.

I then generated several features. I created a geographical region feature to more easily capture similarities across nearby states. I used the information from wikipedia to generate normalized engine and drivetrain labels. The engine is a particularly important feature for Jeep trim levels as certain engines are only available with a subset of trim packages. I generated a feature that extracts a trim label from the unstructured text. For the most of the features that were multi-token strings I also created bag-of-words and bag-of-bigram features. I dropped the words and bigrams that did not occur in more than at least 5 entries. The vehicle history was a standardized list so I just extracted the list of elements.

I then evaluated several different regressor and classifier models. A random forest performed very well for both the trim classification task and the price prediction. The hyperparameters for these models were then chosen by choosing the simplest model that had a cross validation score no more than one standard error below the best model. Each model was trained on the full dataset minus the entries where the predicted feature was null. The predictions are contained in "predictions.csv"

1 Appendix 1: Minimal Code to Generate Predictions

```
#!/usr/bin/env python
# coding: utf-8

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from nltk.metrics import edit_distance
import re, nltk
nltk.download('punkt')

train = pd.read_csv('Training_DataSet.csv')
test = pd.read_csv('Test_Dataset.csv')

# remove capitalization
for x in ["SellerCity", "VehColorExt", "VehColorInt", "VehFeats", "VehSellerNotes"]:
    train[x]=train[x].str.lower()
    test[x]=test[x].str.lower()

train.at[1125, "SellerZip"]=23294
train.at[1125, "SellerListSrc"]="HomeNet Automotive" # used this source for future sales
train.at[1125, "VehFuel"]="Gasoline"

train.at[3855, "SellerZip"]=48124
train.at[3855, "SellerListSrc"]="Inventory Command Center"
```

```

train.at[3855,"VehFuel"]="Gasoline"

# Two records are missing VehMileage and VehListdays, we will just fill those in with the means.

train["VehMileage"].fillna(value=train["VehMileage"].mean(),inplace=True)
test["VehMileage"].fillna(value=test["VehMileage"].mean(),inplace=True)
train["VehListdays"].fillna(value=train["VehListdays"].mean(),inplace=True)

# For the null VehColorExt VehColorInt VehDriveTrain VehEngine VehPriceLabel VehTransmission add an "unknown" category.

for x in ["VehColorExt", "VehColorInt", "VehDriveTrain", "VehEngine", "VehPriceLabel", "VehTransmission"]:
    train[x].fillna(value="unknown",inplace=True)
    test[x].fillna(value="unknown",inplace=True)

# For many names the first token is most informative with the remaining tokens just specifying the location.
# To capture sellers that have multiple locations, add a feature for the first token in a name.

train["SellerShortName"]=train["SellerName"].str.lower().str.split(n=1,expand=True)[0]
test["SellerShortName"]=test["SellerName"].str.lower().str.split(n=1,expand=True)[0]

# There is just one data point for HI and it doesn't show up in testing. HI is likely an outlier so it is removed.

train = train[train["SellerState"]!= "HI"]

# We add a feature that consolidates states in their Census Bureau-designated divisions.

divisions = {"New England" : {"CT", "ME", "MA", "NH", "RI", "VT"},
"Middle Atlantic" : {"NJ", "NY", "PA"},
"South Atlantic" : {"MD", "DC", "DE", "VA", "WV", "NC", "SC", "GA", "FL"},
"East North Central" : {"IL", "IN", "OH", "MI", "WI"},
"East South Central" : {"KY", "TN", "MS", "AL"},
"West North Central" : {"ND", "SD", "NE", "KS", "MO", "IA", "MN"},
"West South Central" : {"TX", "LA", "AR", "OK"},
"Mountain" : {"MT", "ID", "WY", "NV", "UT", "CO", "AZ", "NM"},
"Pacific" : {"AK", "HI", "WA", "OR", "CA"}}
state_to_division = dict([(state,division) for division in divisions.keys() for state in divisions[division]])

train["SellerDivision"] = train["SellerState"].apply(lambda state: state_to_division[state])
test["SellerDivision"] = test["SellerState"].apply(lambda state: state_to_division[state])

# Normalize the labels for 4WD and FWD drivetrains.

drivetrains = {
"4WD" : ["4WD", "AWD", "4X4", "Four Wheel Drive", "ALL-WHEEL DRIVE", "All Wheel Drive",
"4x4", "4x4/4-wheel drive", "4x4/4WD", "AWD or 4x4", "All-wheel Drive", "ALL WHEEL",
"AllWheelDrive", "ALL-WHEEL DRIVE WITH LOCKING AND LIMITED-SLIP DIFFERENTIAL", "4WD/AWD"],
"FWD" : ["FWD", "FRONT-WHEEL DRIVE", "Front Wheel Drive", "Front-wheel Drive", "2WD"],
"unknown" : ["unknown"]}
}
clean_drivetrain = dict([(raw,clean) for clean in drivetrains.keys() for raw in drivetrains[clean]])
train["VehDriveTrainClean"] = train["VehDriveTrain"].apply(lambda raw: clean_drivetrain[raw])
test["VehDriveTrainClean"] = test["VehDriveTrain"].apply(lambda raw: clean_drivetrain[raw])

# Tokenize the engine strings
train["VehEngineTokens"] = train["VehEngine"].str.split()
test["VehEngineTokens"] = test["VehEngine"].str.split()

# Extract the list of features from VehFeats and generate words and bigrams
nonalphanum = re.compile("[^a-zA-Z0-9_ ]")
def words_and_bigrams(s):
words = list(nltk.word_tokenize(s))
bigrams = list(nltk.bigrams(words))
return [(w,) for w in words] + bigrams

def extract_feats(feats):
if pd.isna(feats):
return [("unknown",)]
else:
return [gram for feat in feats.strip()[1:].split(' ', ' ') for gram in words_and_bigrams(nonalphanum.sub(' ', feat).strip())]
train["VehFeatTokens"] = train["VehFeats"].apply(extract_feats)
test["VehFeatTokens"] = test["VehFeats"].apply(extract_feats)

# Extract words and bigrams from Seller notes
def extract_notes(notes):
if pd.isna(notes):
return [("unknown",)]
else:
return [gram for gram in words_and_bigrams(nonalphanum.sub(' ', notes).strip())]
train["VehSellerNotesTokens"] = train["VehSellerNotes"].apply(extract_notes)

```

```

test["VehSellerNotesTokens"] = test["VehSellerNotes"].apply(extract_notes)

# As VehHistory is standardized we just extract the elements in the list
train["VehHistoryTokens"] = train["VehHistory"].fillna('unknown').str.split(' ')
test["VehHistoryTokens"] = test["VehHistory"].fillna('unknown').str.split(' ')

# Extract words and bigrams from seller names
def extract_names(name):
    if pd.isna(name):
        return [("unknown",)]
    else:
        return [gram for gram in words_and_bigrams(nonalphanumeric.sub(' ', name).strip())]
train["SellerNameTokens"] = train["SellerName"].apply(extract_names)
test["SellerNameTokens"] = test["SellerName"].apply(extract_names)

def extract_engine(row):
    if (not pd.isna(row["VehFeats"])) and "3.6" in row["VehFeats"] or (not pd.isna(row["VehSellerNotes"])) and "3.6" in row["VehSellerNotes"]:
        return "3.6L V6"
    elif (not pd.isna(row["VehFeats"])) and "5.7" in row["VehFeats"] or (not pd.isna(row["VehSellerNotes"])) and "5.7" in row["VehSellerNotes"]:
        return "5.7L V8"
    elif (not pd.isna(row["VehFeats"])) and "6.2" in row["VehFeats"] or (not pd.isna(row["VehSellerNotes"])) and "6.2" in row["VehSellerNotes"]:
        return "6.2L V8"
    elif (not pd.isna(row["VehFeats"])) and "6.4" in row["VehFeats"] or (not pd.isna(row["VehSellerNotes"])) and "6.4" in row["VehSellerNotes"]:
        return "6.4L V8"
    else:
        return row["VehEngine"]
def clean_train_engine(row):
    if row["VehMake"] == "Cadillac":
        return "3.6L V6"
    elif row["VehFuel"]=="Diesel":
        return "3.0L V6"
    elif "3.6" in row["VehEngine"] or "V6" in row["VehEngine"] or row["VehEngine"] in ["6-cylinder", "6", "6 Cylinder", "V-6 cyl"] \
    or row["Vehicle_Trim"] in ["Laredo", "Laredo E"]:
        return "3.6L V6"
    elif "6.4" in row["VehEngine"] or row["Vehicle_Trim"] in ["SRT", "SRT8"]:
        return "6.4L V8"
    elif "6.2" in row["VehEngine"] or "Trackhawk" == row["Vehicle_Trim"]:
        return "6.2L V8"
    elif "5.7" in row["VehEngine"] or ("8" in row["VehEngine"] and not (row["Vehicle_Trim"] in ["SRT", "SRT8", "Trackhawk"])):
        return "5.7L V8"
    else:
        return extract_engine(row)
def clean_test_engine(row):
    if row["VehMake"] == "Cadillac":
        return "3.6L V6"
    elif row["VehFuel"]=="Diesel":
        return "3.0L V6"
    elif "3.6" in row["VehEngine"] or "V6" in row["VehEngine"] or row["VehEngine"] in ["6-cylinder", "6", "6 Cylinder", "V-6 cyl"]:
        return "3.6L V6"
    elif "6.2" in row["VehEngine"]:
        return "6.2L V8"
    elif "6.4" in row["VehEngine"] or row["VehEngine"]=="8-cylinder": # the one "8-cylinder" has SRT in notes
        return "6.4L V8"
    elif "5.7" in row["VehEngine"] or row["VehEngine"]=="8 Cylinder Engine":
        # the one "8 Cylinder Engine" has uninformative features, rather than label it "unknown" I label it the most common V8 engine type
        return "5.7L V8"
    elif row["VehEngine"]=="0":
        return "unknown"
    else:
        return extract_engine(row)

train["VehEngineClean"] = train.apply(clean_train_engine,axis=1)
test["VehEngineClean"] = test.apply(clean_test_engine,axis=1)

# A simple attempt to extract a trim from VehFeats or VehSellerNotes
def extract_trim(row):
    if not pd.isnull(row["VehSellerNotes"]):
        if row["VehMake"]=="Cadillac":
            if "platinum" in row["VehSellerNotes"]:
                return "Platinum"
            elif "premium luxury" in row["VehSellerNotes"]:
                return "Premium Luxury"
            elif "luxury" in row["VehSellerNotes"]:
                return "Luxury"
            elif "base" in row["VehSellerNotes"]:
                return "Base"
            else:
                if "laredo e" in row["VehSellerNotes"]:
                    return "Laredo E"
                elif "laredo" in row["VehSellerNotes"]:
                    return "Laredo E"
                elif "overland" in row["VehSellerNotes"]:
                    return "Overland"
                elif "high altitude" in row["VehSellerNotes"]:
                    return "High Altitude"
                elif "altitude" in row["VehSellerNotes"]:
                    return "Altitude"
                elif "summit" in row["VehSellerNotes"]:
                    return "Summit"
                elif "trailhawk" in row["VehSellerNotes"]:
                    return "Trailhawk"
                elif "srt" in row["VehSellerNotes"]:

```

```

return "SRT"
elif "trackhawk" in row["VehSellerNotes"]:
return "Trackhawk"
elif "sterling edition" in row["VehSellerNotes"]:
return "Sterling Edition"
elif "upland" in row["VehSellerNotes"]:
return "Upland"
elif "limited" in row["VehSellerNotes"]:
return "Limited"
elif not pd.isnull(row["VehFeats"]):
if row["VehMake"]=="Cadillac":
if "platinum" in row["VehFeats"]:
return "Platinum"
elif "premium luxury" in row["VehFeats"]:
return "Premium Luxury"
elif "luxury" in row["VehFeats"]:
return "Luxury"
elif "base" in row["VehFeats"]:
return "Base"
else:
if "laredo e" in row["VehFeats"]:
return "Laredo E"
elif "laredo" in row["VehFeats"]:
return "Laredo E"
elif "overland" in row["VehFeats"]:
return "Overland"
elif "high altitude" in row["VehFeats"]:
return "High Altitude"
elif "altitude" in row["VehFeats"]:
return "Altitude"
elif "summit" in row["VehFeats"]:
return "Summit"
elif "trailhawk" in row["VehFeats"]:
return "Trailhawk"
elif "srt" in row["VehFeats"]:
return "SRT"
elif "trackhawk" in row["VehFeats"]:
return "Trackhawk"
elif "sterling edition" in row["VehFeats"]:
return "Sterling Edition"
elif "upland" in row["VehFeats"]:
return "Upland"
elif "limited" in row["VehFeats"]:
return "Limited"
return "Unknown"
train["Vehicle_Trim_Extraction"] = train.apply(extract_trim,axis=1)
test["Vehicle_Trim_Extraction"] = test.apply(extract_trim,axis=1)

# An aggressively normalized version of trim labels
def clean_trim(row):
if pd.isnull(row["Vehicle_Trim"]):
return row["Vehicle_Trim"]
elif row["VehMake"]=="Cadillac":
if row["Vehicle_Trim"] in ["Premium Luxury","Premium Luxury AWD","Premium Luxury FWD"]:
return "Premium Luxury"
elif row["Vehicle_Trim"] in ["Luxury","Luxury AWD","Luxury FWD"]:
return "Luxury"
elif row["Vehicle_Trim"] in ["Platinum","Platinum AWD"]:
return "Platinum"
else:
return "Base"
else:
if row["Vehicle_Trim"] in ["Limited","75th Anniversary","Limited 75th Anniversary Edition","Limited 4x4","75th Anniversary Edition", \
"Limited X","Limited 75th Anniversary"]:
return "Limited"
elif row["Vehicle_Trim"] in ["SRT", "SRT Night"]:
return "SRT"
else:
return row["Vehicle_Trim"]
train["Vehicle_Trim_Clean"] = train.apply(clean_trim,axis=1)

# Replace trim levels that are used only once
train["Vehicle_Trim"].value_counts()
def clean_trim(row):
if pd.isnull(row["Vehicle_Trim"]):
return row["Vehicle_Trim"]
elif row["VehMake"]=="Cadillac":
return row["Vehicle_Trim"]
else:
if row["Vehicle_Trim"] in ["Limited X", "Limited 4x4"]:
return "Limited"
elif row["Vehicle_Trim"] in ["Limited 75th Anniversary","75th Anniversary Edition"]:
return "75th Anniversary"
else:
return row["Vehicle_Trim"]
train["Vehicle_Trim_Semiclean"] = train.apply(clean_trim,axis=1)

# normalize unknown color labeld
def clean_ext_color(color):
if color == "undetermined" or color == "unspecified":
return "unknown"
else:

```

```

return color
train["VehColorExt"] = train["VehColorExt"].apply(clean_ext_color)
test["VehColorExt"] = test["VehColorExt"].apply(clean_ext_color)

# Extract words and bigrams from colors
def extract_color(color):
    if pd.isna(color):
        return [{"unknown",}]
    else:
        return [gram for gram in words_and_bigrams(nonalphaum.sub(' ', color).strip())]
train["VehColorExtTokens"] = train["VehColorExt"].apply(extract_color)
test["VehColorExtTokens"] = test["VehColorExt"].apply(extract_color)
train["VehColorIntTokens"] = train["VehColorInt"].apply(extract_color)
test["VehColorIntTokens"] = test["VehColorInt"].apply(extract_color)

# Encode features for model training
from sklearn import preprocessing
from sklearn.preprocessing import OneHotEncoder, MultiLabelBinarizer, StandardScaler
real_features = ["SellerRating", "SellerRevCnt", "VehListdays", "VehMileage"]
already_encoded_features = ["SellerIsPriv", "VehCertified"]
one_hot_encoded_features = ["SellerListSrc", "SellerState", "SellerDivision", "VehDriveTrainClean", "VehEngineClean", "VehFuel", \
    "VehMake", "VehPriceLabel", "VehYear", "Vehicle_Trim_Extraction"]
mlb_encoded_features = ["SellerNameTokens", "VehEngineTokens", "VehFeatTokens", "VehSellerNotesTokens", "VehHistoryTokens", \
    "VehColorExtTokens", "VehColorIntTokens"]

# Scale real valued features
scaler = StandardScaler().fit(pd.concat([train[real_features], test[real_features]], ignore_index=True))
X_train = pd.DataFrame(scaler.transform(train[real_features]), columns = real_features, index=train.index)
X_predict = pd.DataFrame(scaler.transform(test[real_features]), columns = real_features, index=test.index)

# One hot encode categorical variables
def custom_combiner(feature, category):
    return str(feature) + "_" + str(category)
ohe = OneHotEncoder(feature_name_combiner=custom_combiner, handle_unknown='ignore', sparse_output=False).fit(train[one_hot_encoded_features])
temp = pd.DataFrame(ohe.transform(train[one_hot_encoded_features]), index=train.index, columns=ohe.get_feature_names_out())
X_train = X_train.join(temp)

temp_predict = pd.DataFrame(ohe.transform(test[one_hot_encoded_features]), index=test.index, columns=ohe.get_feature_names_out())
X_predict = X_predict.join(temp_predict)

# Encode lists of features using a multi label binarizer
for feature in mlb_encoded_features:
    mlb = MultiLabelBinarizer(sparse_output=True).fit(pd.concat([train[feature], test[feature]], ignore_index=True))
    temp = pd.DataFrame(sparse.from_spmatrix(mlb.transform(train[feature]), columns=[feature+str(c) for c in mlb.classes_], index=train.index)
    temp = temp.loc[:, temp.sum()>=5] # drop features that don't occur in at least 5 samples in training
    X_train = X_train.join(temp)

temp_predict = pd.DataFrame(sparse.from_spmatrix(mlb.transform(test[feature]), columns=[feature+str(c) for c in mlb.classes_], index=test.index)
temp_predict = temp_predict.loc[:, temp_predict.sum()>=5].columns # drop features that don't occur in at least 5 samples in training
X_predict = X_predict.join(temp_predict)

Y_train = train["Vehicle_Trim"]
Y_train_clean = train["Vehicle_Trim_Clean"]
Y_train_semiclean = train["Vehicle_Trim_Semiclean"]

# Dataset for trim prediction
X_train_trim = X_train[~Y_train.isnull()]
Y_train_trim = Y_train[~Y_train.isnull()]
Y_train_clean_trim = Y_train_clean[~Y_train.isnull()]
Y_train_semiclean_trim = Y_train_semiclean[~Y_train.isnull()]

# Dataset for price prediction
X_train_price = X_train[~train["Dealer_Listing_Price"].isnull()]
feature_scaler = StandardScaler().fit(X_train)
X_train_price_normalized = pd.DataFrame(feature_scaler.transform(X_train_price), columns = X_train_price.columns, index=X_train_price.index)
X_predict_normalized = pd.DataFrame(feature_scaler.transform(X_predict), columns = X_predict.columns, index=X_predict.index)

Y_train_price = train[~train["Dealer_Listing_Price"].isnull()]["Dealer_Listing_Price"]
price_scaler = StandardScaler().fit(Y_train_price.to_numpy().reshape(-1, 1))
Y_train_price_normalized = pd.Series(price_scaler.transform(Y_train_price.to_numpy().reshape(-1, 1)).reshape(1,-1)[0], index=Y_train_price.index)

from sklearn.ensemble import RandomForestClassifier
rf_clf = RandomForestClassifier(max_depth=18)
rf_clf.fit(X_train_trim, Y_train_semiclean_trim)
predict = pd.DataFrame(rf_clf.predict(X_predict), columns=["Vehicle_Trim_Predicted"]).join(test["ListingID"])

from sklearn.ensemble import RandomForestRegressor
rf_reg = RandomForestRegressor(n_jobs=-1)
rf_reg.fit(X_train_price_normalized, Y_train_price_normalized)
predict = pd.DataFrame(price_scaler.inverse_transform(rf_reg.predict(X_predict_normalized).reshape(-1, 1)), columns=["Predicted_price"], \
    index=X_predict_normalized.index).join(predict)

predict[["ListingID", "Vehicle_Trim_Predicted", "Predicted_price"]].to_csv("predictions.csv", index=False)

```

2 Appendix 2: Full Jupyter Notebook

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from nltk.metrics import edit_distance
import re, nltk
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to
[nltk_data] /Users/hazeneckert/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
[1]: True
```

```
[2]: train = pd.read_csv('Training_DataSet.csv')
test = pd.read_csv('Test_Dataset.csv')
```

3 Info about the vehicles

Cadillac XT5 2017-2019 - Only offers a 3.6L V6 gasoline engine in the US - Two drivetrains AWD or FWD - One transmission 8-speed - Four Trim levels [Base, Luxury, Premium Luxury, Platinum]

Jeep Grand Cherokee 2015-2019 - 10 trim levels: [Laredo, Laredo E, Altitude, High Altitude, Upland, Trailhawk, Limited, Sterling Edition, Overland, Summit, SRT, SRT Trackhawk] - Offers 6 engine options: - 3.6L V6: Laredo, Laredo E, Altitude, Limited, Sterling Edition, Trailhawk, Overland, High Altitude, Summit - 5.7L V8: Limited, Limited X, Sterling Edition, Trailhawk, Overland, High Altitude, Summit - 6.4L V8: SRT - 6.2L V8: SRT Trackhawk - 3.0L Diesel: Limited, Overland, Summit - The transmission is a function of year and engine - No info about drivetrain

4 Exploratory Data Analysis

ListingID is a unique identifier.

SellerCity - has some misspelled cities: 'Morganton' and 'Morgantown', 'Milwaukee' and 'Milwaukie', - has different capitalizations: 'Green Bay' and 'Green bay', - has some unicode symbols: 'Coeur d'Alene' and 'O'Fallon', - has different labels for the same location: 'Charter Twp of Clinton', 'Clinton','Clinton Township' - is never null in training or testing - There are cities in the test data never observed in the training data

SellerIsPriv is extremely biased towards False, 6284/6298 and is never null in training or testing.

SellerListSrc takes 8 non-null values and is null for 2 entries. These 2 null entries are also the null entries for SellerZip. All values observed in the test set is in the training set. There are no null values in the training data.

SellerName is never null. There are sellers in the test data not in the training data. There are many entries that are very similar like 'CarMax Buffalo' and 'CarMax Buford'.

SellerRating is never null in training or testing. There are 5988 nonzero entries.

SellerRevCnt is never null in training or testing. There are 6128 nonzero entries.

SellerState has observations for each state in training. 6 states are not included in the testing data, ['HI', 'ME', 'MT', 'NM', 'OR', 'VT']. Never null in training or testing. Probably a good idea to remove HI

VehBodystyle is always 'SUV'

VehCertified is mostly false, 4879/6298. Never null in training or testing.

VehColorExt has null values in training and testing but most entries are in the top 20 or so most common colors.

VehColorInt has null values in training and testing but most entries are in the top 10 or so most common colors. Maybe extracting the leather option would be helpful.

VehDriveTrain has null values in training and testing. There are several different names for the same thing.

VehEngine has null values in training and testing. There are several different names for the same thing.

Skip VehFeats for now

VehFuel has null values in the training but not the test data. Other than 'Unknown' and null entries this column is clean.

VehHistory has null values in training and testing. It seems easy to extract a small number of features from this column.

VehListdays is null for the same two entries that SellerListSrc is null. It is never null in the test set.

VehMake is never null in training or testing. It is either 'Jeep' or 'Cadillac'

VehMileage is null in training and test but very rarely, twice and once, respectively.

VehModel is never null in training or testing. It is either 'Grand Cherokee' or 'XT5'

VehPriceLabel is null in training and test but otherwise takes 3 values.

Skip VehSellerNotes for now

VehType is always 'Used'

VehTransmission is null in training and test and has a bunch of redundant labels.

VehYear is never null and the training and testing share the same 5 years.

Vehicle_Trim is null for 405 entries in training.

Dealer_Listing_Price is null for 52 entries in training.

```
[3]: train.describe(include='all')
```

```
[3]:
```

	ListingID	SellerCity	SellerIsPriv	SellerListSrc	\
count	6.298000e+03	6298	6298	6296	
unique	NaN	1318	2	8	
top	NaN	Chicago	False	Digital Motorworks (DMi)	
freq	NaN	118	6284	3086	
mean	4.318130e+06	NaN	NaN	NaN	
std	2.486031e+06	NaN	NaN	NaN	
min	3.287000e+03	NaN	NaN	NaN	
25%	2.178112e+06	NaN	NaN	NaN	
50%	4.298122e+06	NaN	NaN	NaN	
75%	6.488249e+06	NaN	NaN	NaN	
max	8.620012e+06	NaN	NaN	NaN	

	SellerName	SellerRating	\
count	6298	6298.000000	
unique	2452	NaN	
top	Vroom (Online Dealer - Nationwide Delivery)	NaN	
freq	381	NaN	
mean	NaN	4.138346	
std	NaN	1.188033	
min	NaN	0.000000	
25%	NaN	4.000000	
50%	NaN	4.600000	
75%	NaN	4.800000	
max	NaN	5.000000	

	SellerRevCnt	SellerState	SellerZip	VehBodystyle	...	VehMake	\
count	6298.000000	6298	6296.000000	6298	...	6298	
unique	NaN	50	NaN	1	...	2	
top	NaN	IL	NaN	SUV	...	Jeep	
freq	NaN	753	NaN	6298	...	4199	
mean	434.565576	NaN	45234.211722	NaN	...	NaN	
std	1274.257411	NaN	20380.478191	NaN	...	NaN	
min	0.000000	NaN	1105.000000	NaN	...	NaN	
25%	28.000000	NaN	28806.000000	NaN	...	NaN	
50%	126.000000	NaN	46410.000000	NaN	...	NaN	
75%	401.000000	NaN	60126.000000	NaN	...	NaN	
max	14635.000000	NaN	99654.000000	NaN	...	NaN	

	VehMileage	VehModel	VehPriceLabel	\
count	6296.000000	6298	6013	
unique	NaN	2	3	
top	NaN	Grand Cherokee	Good Deal	
freq	NaN	4199	4488	
mean	26369.364358	NaN	NaN	
std	13036.568712	NaN	NaN	
min	0.000000	NaN	NaN	
25%	16835.000000	NaN	NaN	
50%	26181.000000	NaN	NaN	
75%	36468.500000	NaN	NaN	
max	83037.000000	NaN	NaN	

	VehSellerNotes	VehType	\
count		6055	6298
unique		4920	1
top	CARVANA CERTIFIED INCLUDES: 150-POINT INSPECTI...	Used	
freq		218	6298
mean		NaN	NaN
std		NaN	NaN
min		NaN	NaN
25%		NaN	NaN
50%		NaN	NaN
75%		NaN	NaN
max		NaN	NaN

	VehTransmission	VehYear	Vehicle_Trim	Dealer_Listing_Price
count	6101	6298.000000	5893	6246.000000
unique	33	NaN	29	NaN
top	8-Speed Automatic	NaN	Limited	NaN
freq	4395	NaN	1912	NaN
mean	NaN	2016.792633	NaN	32265.053314
std	NaN	1.206566	NaN	7538.339005
min	NaN	2015.000000	NaN	18289.000000
25%	NaN	2015.000000	NaN	26900.000000
50%	NaN	2017.000000	NaN	31455.500000
75%	NaN	2018.000000	NaN	35991.000000
max	NaN	2019.000000	NaN	89500.000000

[11 rows x 29 columns]

```
[4]: train.isnull().sum(),test.isnull().sum()
```

```
[4]: (ListingID          0
      SellerCity        0
      SellerIsPriv      0
      SellerListSrc      2
      SellerName         0
      SellerRating       0
      SellerRevCnt       0
      SellerState        0
      SellerZip          2
      VehBodystyle       0
      VehCertified       0
      VehColorExt        73
      VehColorInt       728
      VehDriveTrain     401
      VehEngine          361
      VehFeats           275
      VehFuel            2
      VehHistory        201
      VehListdays       2
      VehMake            0
      VehMileage         2
      VehModel           0
      VehPriceLabel     285
      VehSellerNotes    243
      VehType            0
      VehTransmission   197
      VehYear            0
      Vehicle_Trim      405
      Dealer_Listing_Price 52
      dtype: int64,
      ListingID          0
      SellerCity        0
      SellerIsPriv      0
      SellerListSrc      0
      SellerName         0
      SellerRating       0
      SellerRevCnt       0
      SellerState        0
      SellerZip          0
      VehBodystyle       0
      VehCertified       0
      VehColorExt         7
      VehColorInt       108
      VehDriveTrain      64
      VehEngine          58
      VehFeats           37
      VehFuel            0
      VehHistory         27
      VehListdays       0
      VehMake            0
      VehMileage         1
      VehModel           0
      VehPriceLabel      38
      VehSellerNotes     41
      VehType            0
      VehTransmission    27
      VehYear            0
      dtype: int64)
```

```
[5]: # remove capitalization
      for x in ["SellerCity","VehColorExt","VehColorInt","VehFeats", "VehSellerNotes"]:
          train[x]=train[x].str.lower()
          test[x]=test[x].str.lower()
```

```
[6]: train.iloc[3855]
```

```
[6]: ListingID          5306897
      SellerCity        dearborn
      SellerIsPriv      False
      SellerListSrc      NaN
      SellerName         Jack Demmer Lincoln
      SellerRating       4.8
      SellerRevCnt       261
      SellerState        MI
      SellerZip          NaN
      VehBodystyle       SUV
      VehCertified       False
      VehColorExt        NaN
      VehColorInt        NaN
      VehDriveTrain      NaN
      VehEngine          NaN
      VehFeats           NaN
      VehFuel            NaN
      VehHistory         NaN
      VehListdays       NaN
```



```

VehMake           Jeep
VehMileage        36678.0
VehModel          Grand Cherokee
VehPriceLabel     Fair Price
VehSellerNotes    NaN
VehType           Used
VehTransmission   NaN
VehYear           2015
Vehicle_Trim      Limited
Dealer_Listing_Price 23500.0
Name: 3855, dtype: object

```

```
[7]: # Find the entries with null SellerListSrc
train[train.SellerListSrc.isnull()]
```

```

[7]:   ListingID SellerCity SellerIsPriv SellerListSrc \
1125   1562581 richmond      False      NaN
3855   5306897 dearborn      False      NaN

      SellerName SellerRating SellerRevCnt SellerState \
1125 Pearson Chrysler Jeep Dodge RAM      1.0      4      VA
3855 Jack Demmer Lincoln      4.8      261      MI

      SellerZip VehBodystyle ... VehMake VehMileage VehModel \
1125      NaN      SUV ... Jeep 38329.0 Grand Cherokee
3855      NaN      SUV ... Jeep 36678.0 Grand Cherokee

      VehPriceLabel VehSellerNotes VehType VehTransmission VehYear \
1125      Good Deal      NaN      Used      NaN      2017
3855      Fair Price      NaN      Used      NaN      2015

      Vehicle_Trim Dealer_Listing_Price
1125      Limited      26333.0
3855      Limited      23500.0

[2 rows x 29 columns]

```

```
[8]: train[train.SellerName.str.contains("Lincoln")]["SellerListSrc"].value_counts()
```

```

[8]: SellerListSrc
Inventory Command Center    39
Digital Motorworks (DMi)   34
HomeNet Automotive         10
Name: count, dtype: int64

```

Only two records have null SellerListSrc, VehFuel and SellerZip, records 1125 and 3855. Rather than drop these record or create an ‘unknown’ category, we are going to fill these in the best we can. We can look up the zipcodes for these dealerships. We have other entries from the seller in record 1125, which we will use to complete the SellerListSrc. For record 3855 we fill SellerListSrc using the most common source for Lincoln dealerships. We chose to fill in the VehFuel entries with Gasoline since it is the most common.

```

[9]: train.at[1125,"SellerZip"]=23294
train.at[1125,"SellerListSrc"]="HomeNet Automotive" # used this source for future sales
train.at[1125,"VehFuel"]="Gasoline"

train.at[3855,"SellerZip"]=48124
train.at[3855,"SellerListSrc"]="Inventory Command Center"
train.at[3855,"VehFuel"]="Gasoline"

```

Two records are missing VehMileage and VehListdays, we will just fill those in with the means.

```

[10]: train["VehMileage"].fillna(value=train["VehMileage"].mean(),inplace=True)
train["VehMileage"].fillna(value=test["VehMileage"].mean(),inplace=True)
train["VehListdays"].fillna(value=train["VehListdays"].mean(),inplace=True)

```

For the null VehColorExt VehColorInt VehDriveTrain VehEngine VehPriceLabel VehTransmission add an “unknown” category.

```

[11]: for x in ["VehColorExt", "VehColorInt", "VehDriveTrain", "VehEngine", "VehPriceLabel", "VehTransmission"]:
train[x].fillna(value="unknown",inplace=True)
test[x].fillna(value="unknown",inplace=True)

```

For many names the first token is most informative with the remaining tokens just specifying the location. To capture sellers that have multiple locations, add a feature for the first token in a name.

```

[12]: train["SellerShortName"]=train["SellerName"].str.lower().str.split(n=1,expand=True)[0]
test["SellerShortName"]=test["SellerName"].str.lower().str.split(n=1,expand=True)[0]

```

There is just one data point for HI and it doesn't show up in testing. HI is likely an outlier so it is removed.

```
[13]: train = train[train["SellerState"]!= "HI"]
```

We add a feature that consolidates states in their Census Bureau–designated divisions.

```
[14]: divisions = {"New England" : {"CT", "ME", "MA", "NH", "RI", "VT"},
    "Middle Atlantic" : {"NJ", "NY", "PA"},
    "South Atlantic" : {"MD", "DC", "DE", "VA", "WV", "NC", "SC", "GA", "FL"},
    "East North Central" : {"IL", "IN", "OH", "MI", "WI"},
    "East South Central" : {"KY", "TN", "MS", "AL"},
    "West North Central" : {"ND", "SD", "NE", "KS", "MO", "IA", "MN"},
    "West South Central" : {"TX", "LA", "AR", "OK"},
    "Mountain" : {"MT", "ID", "WY", "NV", "UT", "CO", "AZ", "NM"},
    "Pacific" : {"AK", "HI", "WA", "OR", "CA"}}

state_to_division = dict([(state, division) for division in divisions.keys() for state in divisions[division]])

train["SellerDivision"] = train["SellerState"].apply(lambda state: state_to_division[state])
test["SellerDivision"] = test["SellerState"].apply(lambda state: state_to_division[state])
```

Normalize the labels for 4WD and FWD drivetrains.

```
[15]: drivetrains = {
    "4WD" : ["4WD", "AWD", "4X4", "Four Wheel Drive", "ALL-WHEEL DRIVE", "All Wheel Drive",
    "4x4", "4x4/4-wheel drive", "4x4/4WD", "AWD or 4x4", "All-wheel Drive", "ALL WHEEL",
    "AllWheelDrive", "ALL-WHEEL DRIVE WITH LOCKING AND LIMITED-SLIP DIFFERENTIAL", "4WD/AWD"],
    "FWD" : ["FWD", "FRONT-WHEEL DRIVE", "Front Wheel Drive", "Front-wheel Drive", "2WD"],
    "unknown" : ["unknown"]
}

clean_drivetrain = dict([(raw, clean) for clean in drivetrains.keys() for raw in drivetrains[clean]])
train["VehDriveTrainClean"] = train["VehDriveTrain"].apply(lambda raw: clean_drivetrain[raw])
test["VehDriveTrainClean"] = test["VehDriveTrain"].apply(lambda raw: clean_drivetrain[raw])
```

```
[16]: # Tokenize the engine strings
train["VehEngineTokens"] = train["VehEngine"].str.split()
test["VehEngineTokens"] = test["VehEngine"].str.split()
```

```
[17]: # Extract the list of features from VehFeats and generate words and bigrams
nonalphanum = re.compile("[a-zA-Z0-9_"]
def words_and_bigrams(s):
    words = list(nltk.word_tokenize(s))
    bigrams = list(nltk.bigrams(words))
    return [(w,) for w in words] + bigrams

def extract_feats(feats):
    if pd.isna(feats):
        return [("unknown",)]
    else:
        return [gram for feat in feats.strip('"').split('" ' ) for gram in words_and_bigrams(nonalphanum.sub(' ', feat).strip())]
train["VehFeatTokens"] = train["VehFeats"].apply(extract_feats)
test["VehFeatTokens"] = test["VehFeats"].apply(extract_feats)
```

```
[18]: # Extract words and bigrams from Seller notes
def extract_notes(notes):
    if pd.isna(notes):
        return [("unknown",)]
    else:
        return [gram for gram in words_and_bigrams(nonalphanum.sub(' ', notes).strip())]
train["VehSellerNotesTokens"] = train["VehSellerNotes"].apply(extract_notes)
test["VehSellerNotesTokens"] = test["VehSellerNotes"].apply(extract_notes)
```

```
[19]: # As VehHistory is standardized we just extract the elements in the list
train["VehHistoryTokens"] = train["VehHistory"].fillna('unknown').str.split(' ')
test["VehHistoryTokens"] = test["VehHistory"].fillna('unknown').str.split(' ')
```

```
[20]: # Extract words and bigrams from seller names
def extract_names(name):
    if pd.isna(name):
        return [("unknown",)]
    else:
        return [gram for gram in words_and_bigrams(nonalphanum.sub(' ', name).strip())]
train["SellerNameTokens"] = train["SellerName"].apply(extract_names)
test["SellerNameTokens"] = test["SellerName"].apply(extract_names)
```

This next feature cleans the VehEngine feature and attempts to extract it from VehFeats or VehSellerNotes if it is unknown.

```
[21]: def extract_engine(row):
    if (not pd.isna(row["VehFeats"]) and "3.6" in row["VehFeats"]) or (not pd.isna(row["VehSellerNotes"]) and "3.6" in row["VehSellerNotes"]):
        return "3.6L V6"
    elif (not pd.isna(row["VehFeats"]) and "5.7" in row["VehFeats"]) or (not pd.isna(row["VehSellerNotes"]) and "5.7" in row["VehSellerNotes"]):
        return "5.7L V8"
    elif (not pd.isna(row["VehFeats"]) and "6.2" in row["VehFeats"]) or (not pd.isna(row["VehSellerNotes"]) and "6.2" in row["VehSellerNotes"]):
        return "6.2L V8"
    elif (not pd.isna(row["VehFeats"]) and "6.4" in row["VehFeats"]) or (not pd.isna(row["VehSellerNotes"]) and "6.4" in row["VehSellerNotes"]):
        return "6.4L V8"
    else:
        return row["VehEngine"]
```

```

def clean_train_engine(row):
    if row["VehMake"] == "Cadillac":
        return "3.6L V6"
    elif row["VehFuel"] == "Diesel":
        return "3.0L V6"
    elif "3.6" in row["VehEngine"] or "V6" in row["VehEngine"] or row["VehEngine"] in ["6-cylinder", "6", "6 Cylinder", "V-6 cyl"] or
    → row["Vehicle_Trim"] in ["Laredo", "Laredo E"]:
        return "3.6L V6"
    elif "6.4" in row["VehEngine"] or row["Vehicle_Trim"] in ["SRT", "SRT8"]:
        return "6.4L V8"
    elif "6.2" in row["VehEngine"] or "Trackhawk" == row["Vehicle_Trim"]:
        return "6.2L V8"
    elif "5.7" in row["VehEngine"] or ("8" in row["VehEngine"] and not (row["Vehicle_Trim"] in ["SRT", "SRT8", "Trackhawk"])):
        return "5.7L V8"
    else:
        return extract_engine(row)
def clean_test_engine(row):
    if row["VehMake"] == "Cadillac":
        return "3.6L V6"
    elif row["VehFuel"] == "Diesel":
        return "3.0L V6"
    elif "3.6" in row["VehEngine"] or "V6" in row["VehEngine"] or row["VehEngine"] in ["6-cylinder", "6", "6 Cylinder", "V-6 cyl"]:
        return "3.6L V6"
    elif "6.2" in row["VehEngine"]:
        return "6.2L V8"
    elif "6.4" in row["VehEngine"] or row["VehEngine"] == "8-cylinder": # the one "8-cylinder" has SRT in notes
        return "6.4L V8"
    elif "5.7" in row["VehEngine"] or row["VehEngine"] == "8 Cylinder Engine": # the one "8 Cylinder Engine" has uninformative features, rather
    → than label it "unknown" I label it the most common V8 engine type
        return "5.7L V8"
    elif row["VehEngine"] == "0":
        return "unknown"
    else:
        return extract_engine(row)

train["VehEngineClean"] = train.apply(clean_train_engine, axis=1)
test["VehEngineClean"] = test.apply(clean_test_engine, axis=1)

```

[22]: # A simple attempt to extract a trim label from VehFeats or VehSellerNotes

```

def extract_trim(row):
    if not pd.isnull(row["VehSellerNotes"]):
        if row["VehMake"] == "Cadillac":
            if "platinum" in row["VehSellerNotes"]:
                return "Platinum"
            elif "premium luxury" in row["VehSellerNotes"]:
                return "Premium Luxury"
            elif "luxury" in row["VehSellerNotes"]:
                return "Luxury"
            elif "base" in row["VehSellerNotes"]:
                return "Base"
        else:
            if "laredo e" in row["VehSellerNotes"]:
                return "Laredo E"
            elif "laredo" in row["VehSellerNotes"]:
                return "Laredo E"
            elif "overland" in row["VehSellerNotes"]:
                return "Overland"
            elif "high altitude" in row["VehSellerNotes"]:
                return "High Altitude"
            elif "altitude" in row["VehSellerNotes"]:
                return "Altitude"
            elif "summit" in row["VehSellerNotes"]:
                return "Summit"
            elif "trailhawk" in row["VehSellerNotes"]:
                return "Trailhawk"
            elif "srt" in row["VehSellerNotes"]:
                return "SRT"
            elif "trackhawk" in row["VehSellerNotes"]:
                return "Trackhawk"
            elif "sterling edition" in row["VehSellerNotes"]:
                return "Sterling Edition"
            elif "upland" in row["VehSellerNotes"]:
                return "Upland"
            elif "limited" in row["VehSellerNotes"]:
                return "Limited"
    elif not pd.isnull(row["VehFeats"]):
        if row["VehMake"] == "Cadillac":
            if "platinum" in row["VehFeats"]:
                return "Platinum"
            elif "premium luxury" in row["VehFeats"]:
                return "Premium Luxury"
            elif "luxury" in row["VehFeats"]:
                return "Luxury"
            elif "base" in row["VehFeats"]:
                return "Base"
        else:
            if "laredo e" in row["VehFeats"]:
                return "Laredo E"
            elif "laredo" in row["VehFeats"]:
                return "Laredo E"

```

```

        elif "overland" in row["VehFeats"]:
            return "Overland"
        elif "high altitude" in row["VehFeats"]:
            return "High Altitude"
        elif "altitude" in row["VehFeats"]:
            return "Altitude"
        elif "summit" in row["VehFeats"]:
            return "Summit"
        elif "trailhawk" in row["VehFeats"]:
            return "Trailhawk"
        elif "srt" in row["VehFeats"]:
            return "SRT"
        elif "trackhawk" in row["VehFeats"]:
            return "Trackhawk"
        elif "sterling edition" in row["VehFeats"]:
            return "Sterling Edition"
        elif "upland" in row["VehFeats"]:
            return "Upland"
        elif "limited" in row["VehFeats"]:
            return "Limited"
        return "Unknown"
train["Vehicle_Trim_Extraction"] = train.apply(extract_trim,axis=1)
test["Vehicle_Trim_Extraction"] = test.apply(extract_trim,axis=1)

```

```

[23]: # An aggressively normalized version of trim labels
def clean_trim(row):
    if pd.isnull(row["Vehicle_Trim"]):
        return row["Vehicle_Trim"]
    elif row["VehMake"]=="Cadillac":
        if row["Vehicle_Trim"] in ["Premium Luxury","Premium Luxury AWD","Premium Luxury FWD"]:
            return "Premium Luxury"
        elif row["Vehicle_Trim"] in ["Luxury","Luxury AWD","Luxury FWD"]:
            return "Luxury"
        elif row["Vehicle_Trim"] in ["Platinum","Platinum AWD"]:
            return "Platinum"
        else:
            return "Base"
    else:
        if row["Vehicle_Trim"] in ["Limited","75th Anniversary","Limited 75th Anniversary Edition","Limited 4x4","75th Anniversary_
↳Edition","Limited X","Limited 75th Anniversary"]:
            return "Limited"
        elif row["Vehicle_Trim"] in ["SRT", "SRT Night"]:
            return "SRT"
        else:
            return row["Vehicle_Trim"]
train["Vehicle_Trim_Clean"] = train.apply(clean_trim,axis=1)

```

```

[24]: # Replace trim levels that are used only once
train["Vehicle_Trim"].value_counts()
def clean_trim(row):
    if pd.isnull(row["Vehicle_Trim"]):
        return row["Vehicle_Trim"]
    elif row["VehMake"]=="Cadillac":
        return row["Vehicle_Trim"]
    else:
        if row["Vehicle_Trim"] in ["Limited X", "Limited 4x4"]:
            return "Limited"
        elif row["Vehicle_Trim"] in ["Limited 75th Anniversary","75th Anniversary Edition"]:
            return "75th Anniversary"
        else:
            return row["Vehicle_Trim"]
train["Vehicle_Trim_Semiclean"] = train.apply(clean_trim,axis=1)

```

```

[25]: # normalize unknown color labeld
def clean_ext_color(color):
    if color == "undetermined" or color == "unspecified":
        return "unknown"
    else:
        return color
train["VehColorExt"] = train["VehColorExt"].apply(clean_ext_color)
test["VehColorExt"] = test["VehColorExt"].apply(clean_ext_color)

```

```

[26]: # Extract words and bigrams from colors
def extract_color(color):
    if pd.isna(color):
        return [("unknown",)]
    else:
        return [gram for gram in words_and_bigrams(nonalphanum.sub(' ', color).strip())]
train["VehColorExtTokens"] = train["VehColorExt"].apply(extract_color)
test["VehColorExtTokens"] = test["VehColorExt"].apply(extract_color)
train["VehColorIntTokens"] = train["VehColorInt"].apply(extract_color)
test["VehColorIntTokens"] = test["VehColorInt"].apply(extract_color)

```

```

[27]: # Encode features for model training
from sklearn import preprocessing
from sklearn.preprocessing import OneHotEncoder, MultiLabelBinarizer,StandardScaler

```

```

real_features = ["SellerRating", "SellerRevCnt", "VehListdays", "VehMileage"]
already_encoded_features = ["SellerIsPriv", "VehCertified"]
one_hot_encoded_features = []
↪ ["SellerListSrc", "SellerState", "SellerDivision", "VehDriveTrainClean", "VehEngineClean", "VehFuel", "VehMake", "VehPriceLabel", "VehYear", "Vehicle_Trim_Extraction"]
mlb_encoded_features = []
↪ ["SellerNameTokens", "VehEngineTokens", "VehFeatTokens", "VehSellerNotesTokens", "VehHistoryTokens", "VehColorExtTokens", "VehColorIntTokens"]

# Scale real valued features
scaler = StandardScaler().fit(pd.concat([train[real_features], test[real_features]], ignore_index=True))
X_train = pd.DataFrame(scaler.transform(train[real_features]), columns = real_features, index=train.index)
X_predict = pd.DataFrame(scaler.transform(test[real_features]), columns = real_features, index=test.index)

# One hot encode categorical variables
def custom_combiner(feature, category):
    return str(feature) + "_" + str(category)
ohe = OneHotEncoder(feature_name_combiner=custom_combiner, handle_unknown='ignore', sparse_output=False).fit(train[one_hot_encoded_features])
temp = pd.DataFrame(ohe.transform(train[one_hot_encoded_features]), index=train.index, columns=ohe.get_feature_names_out())
X_train = X_train.join(temp)

temp_predict = pd.DataFrame(ohe.transform(test[one_hot_encoded_features]), index=test.index, columns=ohe.get_feature_names_out())
X_predict = X_predict.join(temp_predict)

# Encode lists of features using a multi label binarizer
for feature in mlb_encoded_features:
    mlb = MultiLabelBinarizer(sparse_output=True).fit(pd.concat([train[feature], test[feature]], ignore_index=True))
    temp = pd.DataFrame(sparse.from_spmatrix(mlb.transform(train[feature]), columns=[feature+str(c) for c in mlb.classes_], index=train.index))
    temp = temp.loc[:, temp.sum()>=5] # drop features that don't occur in at least 5 samples in training
    X_train = X_train.join(temp)

    temp_predict = pd.DataFrame(sparse.from_spmatrix(mlb.transform(test[feature]), columns=[feature+str(c) for c in mlb.classes_], index=test.
↪ index))
    temp_predict = temp_predict.loc[:, temp.loc[:, temp.sum()>=5].columns] # drop features that don't occur in at least 5 samples in training
    X_predict = X_predict.join(temp_predict)

Y_train = train["Vehicle_Trim"]
Y_train_clean = train["Vehicle_Trim_Clean"]
Y_train_semiclean = train["Vehicle_Trim_Semiclean"]

```

```

[28]: # Dataset for trim prediction
X_train_trim = X_train[~Y_train.isnull()]
Y_train_trim = Y_train[~Y_train.isnull()]
Y_train_clean_trim = Y_train_clean[~Y_train.isnull()]
Y_train_semiclean_trim = Y_train_semiclean[~Y_train.isnull()]

```

```

[29]: # Dataset for price prediction
X_train_price = X_train[~train["Dealer_Listing_Price"].isnull()]
feature_scaler = StandardScaler().fit(X_train)
X_train_price_normalized = pd.DataFrame(feature_scaler.transform(X_train_price), columns = X_train_price.columns, index=X_train_price.index)
X_predict_normalized = pd.DataFrame(feature_scaler.transform(X_predict), columns = X_predict.columns, index=X_predict.index)

Y_train_price = train[~train["Dealer_Listing_Price"].isnull()]["Dealer_Listing_Price"]
price_scaler = StandardScaler().fit(Y_train_price.to_numpy().reshape(-1, 1))
Y_train_price_normalized = pd.Series(price_scaler.transform(Y_train_price.to_numpy().reshape(-1, 1)).reshape(1,-1)[0], index=Y_train_price.index)

```

```

/Users/hazeneckert/miniconda3/envs/practice/lib/python3.11/site-
packages/sklearn/utils/validation.py:787: UserWarning: pandas.DataFrame with
sparse columns found.It will be converted to a dense numpy array.
    warnings.warn(
/Users/hazeneckert/miniconda3/envs/practice/lib/python3.11/site-
packages/sklearn/utils/validation.py:787: UserWarning: pandas.DataFrame with
sparse columns found.It will be converted to a dense numpy array.
    warnings.warn(
/Users/hazeneckert/miniconda3/envs/practice/lib/python3.11/site-
packages/sklearn/utils/validation.py:787: UserWarning: pandas.DataFrame with
sparse columns found.It will be converted to a dense numpy array.
    warnings.warn(

```

5 Model Training and Prediciton

See the section on model selection to see why these models and parameters were chosen.

```

[31]: from sklearn.ensemble import RandomForestClassifier
rf_clf = RandomForestClassifier(max_depth=18)
rf_clf.fit(X_train_trim, Y_train_semiclean_trim)
predict = pd.DataFrame(rf_clf.predict(X_predict), columns=["Vehicle_Trim_Predicted"]).join(test["ListingID"])

```

```

/Users/hazeneckert/miniconda3/envs/practice/lib/python3.11/site-
packages/sklearn/utils/validation.py:787: UserWarning: pandas.DataFrame with
sparse columns found.It will be converted to a dense numpy array.
    warnings.warn(
/Users/hazeneckert/miniconda3/envs/practice/lib/python3.11/site-
packages/sklearn/utils/validation.py:787: UserWarning: pandas.DataFrame with
sparse columns found.It will be converted to a dense numpy array.
    warnings.warn(

```

```
[32]: from sklearn.ensemble import RandomForestRegressor
rf_reg = RandomForestRegressor(n_jobs=-1)
rf_reg.fit(X_train_price_normalized, Y_train_price_normalized)
predict = pd.DataFrame(price_scaler.inverse_transform(rf_reg.predict(X_predict_normalized).reshape(-1, 1)), columns=["Predicted_price"],
↳ index=X_predict_normalized.index).join(predict)

[33]: predict[["ListingID", "Vehicle_Trim_Predicted", "Predicted_price"]].to_csv("predictions.csv", index=False)
```

6 Model Selection and Evaluation

I first quickly evaluate the performance of a collection of models on the data. I select the most promising one for each task and find hyperparameters for the simplest model within a std dev of the optimal score.

```
[ ]: from sklearn.model_selection import KFold
from sklearn.metrics import r2_score
from sklearn.ensemble import AdaBoostRegressor, RandomForestRegressor, GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import Lasso
regressors = [AdaBoostRegressor(),
              RandomForestRegressor(n_jobs=-1),
              GradientBoostingRegressor(),
              KNeighborsRegressor(3, n_jobs=-1),
              Lasso(alpha=0.1)]

train_ind, test_ind = next(KFold(n_splits=5).split(X_train_price, Y_train_price_normalized))
for reg in regressors:
    reg.fit(X_train_price_normalized.iloc[train_ind], Y_train_price_normalized.iloc[train_ind])

    print("-"*50)
    print(reg.__class__.__name__)

    print('****Results****')
    r2 = r2_score(Y_train_price_normalized.iloc[test_ind], reg.predict(X_train_price_normalized.iloc[test_ind]))
    print("R2: {:.4%}".format(r2))
```

RandomForestRegressor performs the best.

```
[ ]: from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import RocCurveDisplay, roc_curve, roc_auc_score
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier, GradientBoostingClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
classifiers = [
    KNeighborsClassifier(3),
    RandomForestClassifier(),
    AdaBoostClassifier(),
    GradientBoostingClassifier(),
    GaussianNB(),
    LinearDiscriminantAnalysis(),
    QuadraticDiscriminantAnalysis()]

train_ind, test_ind = next(StratifiedKFold(n_splits=3).split(X_train_trim, Y_train_clean_trim))

for clf in classifiers:
    clf.fit(X_train_trim.iloc[train_ind], Y_train_clean_trim.iloc[train_ind])

    print("-"*50)
    print(clf.__class__.__name__)

    print('****Results****')
    auc = roc_auc_score(Y_train_clean_trim.iloc[test_ind], clf.predict_proba(X_train_trim.iloc[test_ind]), multi_class='ovr')
    print("AUC: {:.4%}".format(auc))
```

The default RandomForestClassifier works extremely well and is very quick to train so I will use a tuned version of it for the trim classifier.

```
[ ]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
parameters = {'criterion': ("gini", "entropy", "log_loss"),
              'max_depth': [5, 10, 25, 50]}
rf = RandomForestClassifier()
clf = GridSearchCV(rf, parameters, cv=3, verbose=3, scoring='roc_auc_ovr', n_jobs=-1)
res = clf.fit(X_train_trim, Y_train_semiclean_trim)
res.cv_results_
```

The criterion doesn't affect performance much and the best depth is 10. I am going to select the gini criterion and explore depth more.

```
[ ]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
parameters = {'criterion': ["gini"],
```

```
        'max_depth' : [8, 12, 15, 18, 20, 22]  
    }  
clf = GridSearchCV(rf, parameters, cv=3, verbose=3, scoring='roc_auc_ovr', n_jobs=-1)  
res = clf.fit(X_train_trim, Y_train_semiclean_trim)  
res.cv_results_
```

The score is maximized at 20 but 18 is within a std dev so it will be chosen.